

Theory Problems: Batch 2

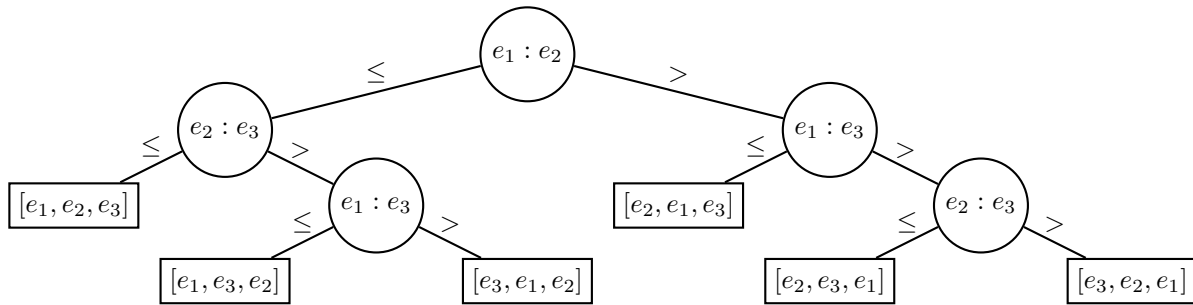
Course: Coursera Algorithms Specialization
Name: Bryan Yaggi

Problem 1

Prove that the worst-case expected running time of every randomized comparison-based sorting algorithm is $\Omega(n \log n)$. (Here the worst-case is over inputs, and the expectation is over the random coin flips made by the algorithm.)

In lecture, we showed that any deterministic comparison-based sorting algorithm has a lower bound, $\Omega(n \log n)$. First, prove the average-case time is $\Omega(n \log n)$. A comparison-based sorting algorithm can be represented as a binary decision tree. See the following example.

Example: Given three elements in a list, $[e_1, e_2, e_3]$.



For n elements, there are $n!$ possible permutations, and each permutation has a $\frac{1}{n!}$ probability of being reached. Let $D(T)$ be the external path length of the decision tree, T , or the sum of the depths of each leaf in the tree. For any T with $k > 1$ leaves, $D(T) = D(LT) + D(RT) + k$, where LT and RT are the left and right subtrees, respectively. This is true because the leaves of LT and RT are one node shallower than in T , and the path to each leaf must go through the root node. Using this observation, the combination of leaves in LT and RT that minimizes $D(T)$ can be found. $D(T) \geq i \log i + (k - i) \log(k - i) + k$ where i is the number of leaves in LT .

$$D(T) \geq i \log i + (k - i) \log(k - i) + k$$

$$\frac{\partial D(T)}{\partial i} = \log i + 1 - \log(k - i) - 1 = \log \frac{i}{k - i}$$

$$\frac{\partial D(T)}{\partial i} = 0 \iff \frac{i}{k - i} = 1 \implies i = \frac{k}{2}$$

This result indicates that $D(T)$ is minimized by a balanced tree. Therefore, $D(T) > k \log k \iff D(T) > n! \log(n!)$. The average-case time, or leaf depth, for sorting n elements is $\log(n!) \implies \Omega(n \log n)$. Now, show the expected running time for random comparison-based sort. A random comparison sort is just a randomly-selected decision tree, T from the set valid decision trees for sorting the input. Let N be the number of valid decision trees.

$$E(\text{time}) = \sum_{j=1}^N p_j \text{time} = \sum_{j=1}^N \frac{1}{N} \Omega(n \log n)$$

$$= \Omega(n \log n)$$

See CLRS Problem 8-1 and solution.

Problem 2

Suppose we modify the deterministic linear-time selection algorithm by grouping the elements into groups of 7, rather than groups of 5. (Use the “median-of-medians” as the pivot, as before.) Does the algorithm still run in $O(n)$ time? What if we use groups of 3?

Answer...

Problem 3

Given an array of n distinct (but unsorted) elements x_1, x_2, \dots, x_n with positive weights w_1, w_2, \dots, w_n such that $\sum_{i=1}^n w_i = W$, a weighted median is an element x_k for which the total weight of all elements with value less than x_k (i.e., $\sum_{x_i < x_k} w_i$) is at most $W/2$, and also the total weight of elements with value larger than x_k (i.e., $\sum_{x_i > x_k} w_i$) is at most $W/2$. Observe that there are at most two weighted medians. Show how to compute all weighted medians in $O(n)$ worst-case time.

Answer...

Problem 4

We showed in an optional video lecture that every undirected graph has only polynomially (in the number n of vertices) different minimum cuts. Is this also true for directed graphs? Prove it or give a counterexample.

Answer...

Problem 5

For a parameter $\alpha \geq 1$, an α -minimum cut is one for which the number of crossing edges is at most α times that of a minimum cut. How many α -minimum cuts can an undirected graph have, as a function of α and the number n of vertices? Prove the best upper bound that you can.

Answer...